

---




# **QLEE:** **Symbolic Execution of QEMU** **Devices with KLEE**

---


## Fairy tale

- [Context] Fuzzing is a popular approach for finding bugs in hypervisors
- [Problem] Existing approaches require **a running instance** of QEMU to perform the fuzzing, resulting in high overhead and complexity.
- [Solution] QLEE use the KLEE symbolic execution engine to automatically test the hypervisors' emulated devices.
- [Evaluation] QLEE allowed for fine-grained testing of these emulators and found one bug/issue in QEMU (yet to be reported).

# Background: Hypervisors



# Background: QEMU Device Emulation



---

# Correctness is Critical

## Hypervisor Security

- Breaking Isolation Between VMs
- Denial-of-Service
- Data Corruption

## Accurate Emulation

- Discrepancies with HW
- Incorrect Device Setup in SW
- Coupling To Implementation Details

# Related Work: VDF


Steps:

1. Record Device Accesses
2. Run AFL:
  - a. Mutate the Recordings
  - b. Replay on a QEMU Instance
3. Minimize Test Case

Limitations:

- Instrumentation code is added **manually**
- Replay requires **running QEMU instances**

Henderson, A. et al. 2017. VDF: Targeted Evolutionary Fuzz Testing of Virtual Devices. *Research in Attacks, Intrusions, and Defenses* (Cham, 2017), 3-25.



# Related Work: HYPER-CUBE

Steps:

1. Boot the OS & enumerate devices
2. Interpret bytecode to interact w/ devices
3. Minimize test-cases

Limitations:

- Requires **Running** QEMU instance
- OS needs drivers and device enumeration logic
- Complex setup

Schumilo, S. et al. 2020. HYPER-CUBE: High-Dimensional Hypervisor Fuzzing. *Proceedings 2020 Network and Distributed System Security Symposium* (San Diego, CA, 2020).





Fig. 2: High-level overview of the system architecture of HYPER-CUBE

# Related Work: NYX (is Complex)



Schumilo, S. et al. Nyx: Greybox Hypervisor Fuzzing using Fast Snapshots and Affine Types.


# Related Work: NYX

Steps:

1. Initial Snapshot
2. Generate/Mutate Payload
3. Execute & Record coverage
4. Restore snapshot if needed

Limitations:

- Nested Hypervisor Setup
- Requires a running hypervisor





## Related Work: Common Pain Point

They are all dynamic, i.e., require executing the target

Observations:

- **Running a hypervisor**
  - Overhead
  - Not Needed to Fuzz a Single Device
- Inputs are not constrained to the device operations (VDF & HYPER-CUBE)
  - The recorded initial seed might have bugs (thanks device driver!)
  - The generated ops might not make any sense for the device
  - NYX, the successor of HYPER-CUBE, has a spec.

---



# **QLEE:** **Symbolic Execution of QEMU** **Devices with KLEE**

# KLEE: Symbolic Execution of LLVM IR

Read IR Instruction	Execute (Update State & Constraints)	Check For Errors (at “dangerous ops”)
define i32 @foo(i32 %a, i32 %b, i32 %c) { entry: ;(1) a + b %sum = add i32 %a, %b  ;(2) c / (a + b) %result = sdiv i32 %c, %sum  ;(3) return result ret i32 %result }	PC  1. sum = a + b 2. result = c / sum 3. return result  Constraints  1. [] 2. [sum != 0] 3. [sum != 0]	(set-logic QF_BV)  (declare-const a (_ BitVec 32)) (declare-const b (_ BitVec 32))  (assert (= (bvadd a b) #x00000000))  ⇒ a = 0, b = 0 ⇒ Division by Zero

---

# QLEE Overview

- Specification for Target Device
  - Target device source path
  - Entry-points w/ their arguments
  - Constraints on the arguments
- Harness Generation
- Stub Generation
  - Global Variables
  - External Functions
- Run KLEE on each Entry-point

# QLEE: Specification

Includes:

- Source path
- Variable declarations
- Constraints & Assignments
- Entry-points

Why YAML?


- Anchor (&) and Alias (\*) to deduplicate
- Order-independent

Why Explicit Arguments?

- Function signature has state as void\*

```
target:  
  source: "hw/char/pl011.c"  
  
variables:  
  state: &state  
    type: "struct PL011State"  
    name: "state"  
    pass-by-pointer: true  
  offset: &offset  
    type: "hwaddr"  
    name: "offset"  
  size: &size  
    type: "unsigned"  
    name: "size"  
  value: &value  
    type: "uint64_t"  
    name: "value"  
  
constraints:  
  - name: "state.fbrd"  
    constraint:  
      - "<= 0x3f"  
  - name: "state.ibrd"  
    constraint:  
      - "<= 0xffff"  
  - name: "state.read_pos"  
    constraint:  
      - ">= 0"  
      - "<= 15"  
  
assignments:  
  - name: "state.id"  
    value: "pl011_id_arm"  
  
entrypoint:  
  - name: "pl011_read"  
    args: [*state, *offset, *size]  
  - name: "pl011_write"  
    args: [*state, *offset, *value, *size]  
  - name: "pl011_get_baudrate"  
    args: [*state]  
  - name: "pl011_can_receive"  
    args: [*state]  
  - name: "pl011_loopback_mdmcctrl"  
    args: [*state]  
  - name: "pl011_reset"  
    args:  
      - << : *state  
        cast: "(DeviceState*)"
```


# QLEE: Harness Generation



```
__attribute__((used))
static int main_pl011_read(void)
{
    struct PL011State state;
    hwaddr offset;
    unsigned size;
    klee_make_symbolic(&state, sizeof(state), "state");
    klee_make_symbolic(&offset, sizeof(offset), "offset");
    klee_make_symbolic(&size, sizeof(size), "size");
    klee_assume(state.fbrd <= 0x3f);
    klee_assume(state.ibrd <= 0xffff);
    klee_assume(state.read_pos >= 0);
    klee_assume(state.read_pos <= 15);
    state.id = pl011_id_arm;
    pl011_read(&state, offset, size);
    return 0;
}
```

---

# QLEE: Stubs Generation



## AutoStub:


- C++ and LLVM IR Library
- Stubs
  - Global variables ⇒ initialized to 0
  - External Functions ⇒ return 0
- Types are preserved

## Why use IR as input?

- Do not have to deal with #include<header>

---

# QLEE: Run KLEE



- Run KLEE per Entry-point
- Output saved to  
`<outputdir>/<timestamp>/<entrypoint>`



# Evaluation

QEMU: commit 72b88908d1

## UART Devices

- Simple & Similar
- Serial Communication  
(e.g., Write character to console)

Target	Entry-points
ARM's PL011	<ul style="list-style-type: none"><li>• pl011_read</li><li>• pl011_write</li><li>• pl011_get_baudrate</li><li>• pl011_can_receive</li><li>• pl011_loopback_mdmctrl</li><li>• pl011_reset</li></ul>
“Serial.c” (Intel’s 16550A)	<ul style="list-style-type: none"><li>• serial_ioport_read</li><li>• serial_ioport_write</li></ul>
AVR USART (ATmega, e.g., Arduino)	<ul style="list-style-type: none"><li>• avr_usart_read</li><li>• avr_usart_write</li></ul>



# Evaluation

	Without Constraints		With <sup>†</sup> Constraints	
	True Positives	False Positives	True Positives	False Positives
PL011	1	2	1	1
“Serial.c” 16550A	0	3*	0	1*
AVR USART	0	3	0	0

<sup>†</sup> Not all device constraints were added

\* Manually terminated potentially infinite loop (undecidable)



# Discussion

## PL011 Bug (?)

- Not reported yet
- Array out-of-bounds access

```
case 0x3f8 ... 0x400:  
    r = s->id[(offset - 0xfe0) >> 2];  
    break;
```

- Case range is a gcc extension
- Case x..y is **inclusive** of both x and y

offset == 0x1000

$$\Rightarrow (0x1000 - 0xfe0 >> 2) = 8$$

But s->id is 8 bytes, i.e.  $0 \leq i \leq 7$

## False Positives:

- Trivial Cases
  - Access size == 1
  - Value < 0xff (one byte access)
  - Struct field is null due to missing setup
- Suppressed with constraints

## Limitation:

- Proper struct initialization
- Stub functions are not “accurate”

---

# Bibliography

- Schumilo, S. et al. Nyx: Greybox Hypervisor Fuzzing using Fast Snapshots and Affine Types.
- Henderson, A. et al. 2017. VDF: Targeted Evolutionary Fuzz Testing of Virtual Devices. Research in Attacks, Intrusions, and Defenses (Cham, 2017), 3–25.
- Schumilo, S. et al. 2020. HYPER-CUBE: High-Dimensional Hypervisor Fuzzing. Proceedings 2020 Network and Distributed System Security Symposium (San Diego, CA, 2020).
- LLVM Project. 2003–present. The LLVM Compiler Infrastructure. Available at: <https://llvm.org>.

